



Ensemble application of convolutional neural networks and multiple kernel learning for multimodal sentiment analysis



Soujanya Poria^a, Haiyun Peng^b, Amir Hussain^a, Newton Howard^c, Erik Cambria^{b,*}

^a Department of Computing Science and Mathematics, University of Stirling, UK

^b School of Computer Science and Engineering, Nanyang Technological University, Singapore

^c Computational Neuroscience and Functional Neurosurgery, University of Oxford, UK

ARTICLE INFO

Article history:

Received 29 September 2015

Revised 4 August 2016

Accepted 22 September 2016

Available online 8 February 2017

Keywords:

Multimodal sentiment analysis

Convolutional neural network

Deep learning

Sentiment

Emotion

MKL

ELM

SVM

Classification

ABSTRACT

The advent of the Social Web has enabled anyone with an Internet connection to easily create and share their ideas, opinions and content with millions of other people around the world. In pace with a global deluge of videos from billions of computers, smartphones, tablets, university projectors and security cameras, the amount of multimodal content on the Web has been growing exponentially, and with that comes the need for decoding such information into useful knowledge. In this paper, a multimodal affective data analysis framework is proposed to extract user opinion and emotions from video content. In particular, multiple kernel learning is used to combine visual, audio and textual modalities. The proposed framework outperforms the state-of-the-art model in multimodal sentiment analysis research with a margin of 10–13% and 3–5% accuracy on polarity detection and emotion recognition, respectively. The paper also proposes an extensive study on decision-level fusion.

© 2017 Elsevier B.V. All rights reserved.

1. Introduction

Subjectivity detection and sentiment analysis consist of the automatic identification of the human mind's private states, e.g., opinions, emotions, moods, behaviors and beliefs [1]. In particular, the former focuses on classifying sentiment data as either objective (neutral) or subjective (opinionated), while the latter aims to infer a positive or negative polarity. Hence, in most cases, both tasks are considered binary classification problems.

To date, most of the work on sentiment analysis has been carried out on text data. With a videocamera in every pocket and the rise of social media, people are now making use of videos (e.g., YouTube, Vimeo, VideoLectures), images (e.g., Flickr, Picasa, Facebook) and audio files (e.g., podcasts) to air their opinions on social media platforms. Thus, it has become critical to find new methods for the mining of opinions and sentiments from these diverse modalities. Plenty of research has been carried out in the field of audio-visual emotion recognition. Some work has also been conducted on fusing audio, visual and textual modalities to detect

emotion from videos. However, a unique common framework is still missing for both tasks. There are also very few studies combining textual clues with audio and visual features. This leads to the need for more extensive research on the use of these three channels together. This paper aims to solve the two key research questions given below –

- Is a common framework useful for both multimodal emotion and sentiment analysis?
- Can audio, visual and textual features jointly enhance the performance of unimodal and bimodal emotion and sentiment analysis classifiers?

Studies conducted in the past lacked extensive research [2–4] and very few of them clearly described the extraction of features and fusion of the information extracted from different modalities. In this paper, we discuss the feature extraction process from different modalities in detail and explain how to use such features for multimodal affect analysis. The YouTube dataset originally developed by [5] and the IEMOCAP dataset [6] were used to demonstrate the accuracy of the proposed framework. We used several supervised classifiers for the sentiment classification task: the CLM-Z [7] based method was used for feature extraction from visual modality; the openSMILE toolkit was used to extract various features from audio; and finally, textual features were extracted

* Corresponding author.

E-mail addresses: sp47@cs.stir.ac.uk (S. Poria), peng0065@ntu.edu.sg (H. Peng), ahu@cs.stir.ac.uk (A. Hussain), newton.howard@nds.ox.ac.uk (N. Howard), cambria@ntu.edu.sg (E. Cambria).

using a deep convolutional neural network (CNN). The fusion of these heterogeneous features was carried out by means of multiple kernel learning (MKL) using support vector machine (SVM) as a classifier with different types of kernel.

The rest of the paper is organized as follows: [Section 2](#) proposes motivations behind this work. [Section 3](#) discusses related works on multimodal emotion detection, sentiment analysis, and multimodal fusion. [Section 4](#) describes the used datasets in detail. [Sections 5, 6, and 7](#) explain how visual, audio, and textual data are processed, respectively. [Section 8](#) proposes experimental results. [Section 9](#) proposes a faster version of the framework. Finally, [Section 10](#) concludes the paper.

2. Motivations

The research in this field is rapidly picking up and has attracted the attention of academia and industry alike. Combined with advances in signal processing and AI, this research has led to the development of advanced intelligent systems that intend to detect and process affective information contained in multimodal sources [16]. However, the majority of such state-of-the-art frameworks rely on processing a single modality, i.e., text, audio, or video. Additionally, all of these systems are known to exhibit limitations in terms of meeting robustness, accuracy, and overall performance requirements, which, in turn, greatly restricts the usefulness of such systems in real-world applications.

The aim of multi-sensor data fusion is to increase the accuracy and reliability of estimates [8]. Many applications, such as navigation tools, have already demonstrated the potential of data fusion. These illustrate the importance and feasibility of developing a multimodal framework that could cope with all three sensing modalities – text, audio, and video in human-centric environments. Humans communicate and express their emotions and sentiments through different channels. Textual, audio, and visual modalities are concurrently and cognitively exploited to enable effective extraction of the semantic and affective information conveyed in conversation [9].

People are gradually shifting from text to video to express their opinion about a product or service, as it is now much easier and faster to produce and share them. For the same reasons, potential customers are now more inclined to browse for video reviews of the product they are interested in, rather than looking for lengthy written reviews. Another reason for doing this is that, while reliable written reviews are quite hard to find, it is sufficient to search for the name of the product on YouTube and choose the clips with most views in order to find good video reviews. Finally, videos are generally more reliable than written text as reviewers often reveal their identity by showing their face, which also allows viewers to better decode conveyed emotions [10].

Hence, videos can be an excellent resource for emotion and sentiment analysis but the medium also comes with major challenges which need to be overcome. For example, expressiveness of opinion varies widely from person to person [2]. Some people express their opinions more vocally, some more visually and others rely exclusively on logic and express little emotion. These personal differences can help guide us towards the affect seeking expression. When a person expresses his or her opinions with more vocal modulation, the audio data will often contain most of the clues indicative of an opinion. When a person is highly communicative via facial expressions, most of the data needed for opinion mining may often be determined through facial expression analysis. So, a generic model needs to be developed which can adapt itself for any user and provide a consistent result. Both of our multimodal affect analysis models are trained on robust data containing opinions or narratives from a wide range of users. In this paper, we show that the ensemble application of feature extraction from

different types of data and modalities is able to significantly enhance the performance of multimodal emotion and sentiment approach.

3. Related work

In this section, we discuss related works in multimodal affect detection covering both emotion and sentiment analysis.

3.1. Text based emotion and sentiment analysis

Sentiment analysis systems can be broadly categorized into knowledge-based and statistics-based systems [11]. While the use of knowledge bases was initially more popular for the identification of emotions and polarity in text, sentiment analysis researchers have recently been using statistics-based approaches, with a special focus on supervised statistical methods. For example, Pang et al. [12] compared the performance of different machine learning algorithms on a movie review dataset and obtained 82.90% accuracy, using only a large number of textual features. A recent approach by Socher et al. [13] obtained even better accuracy (85%) on the same dataset using a recursive neural tensor network (RNTN). Yu and Hatzivassiloglou [14] used semantic orientation of words to identify polarity at sentence level. Melville et al. [15] developed a framework that exploits word-class association information for domain-dependent sentiment analysis.

Other unsupervised or knowledge-based approaches to sentiment analysis include Melville et al. [17], who proposed a mathematical model to extract emotional clues from blogs and then used these for sentiment detection; Gangemi et al. [18], who presented an unsupervised frame-based approach to identify opinion holders and topics based on the assumption that events and situations are the primary entities for contextualizing opinions; and Cambria et al. [19], who proposed a multidisciplinary framework for polarity detection based on SenticNet [20], a concept-level common-sense knowledge base.

Sentiment analysis research can also be categorized as single-domain [12,16,21,22] or cross-domain [23]. The work presented in [24] discusses spectral feature alignment to group domain-specific words from different domains into clusters. They first incorporated domain-independent words to help the clustering process and then exploited the resulting clusters to reduce the gap between domain-specific words of two domains. Bollegala et al. [25] developed a sentiment-sensitive distributional thesaurus by using labeled training data from the source domain and unlabeled training data from both the source and target domains. Sentiment sensitivity was obtained by including documents' sentiment labels into the context vector. At the time of training and testing, this sentiment thesaurus was used to expand the feature vector.

The task of automatically identifying fine-grained emotions, such as anger, joy, surprise, fear, disgust, and sadness, explicitly or implicitly expressed in a text has been addressed by several researchers [26,27]. There are a number of theories on emotion taxonomy which spans from Ekman's emotion categorization model to the Hourglass of Emotion (Fig. 1) [28]. So far, approaches to text-based emotion and sentiment detection rely mainly on rule-based techniques, bag of words modeling using a large sentiment or emotion lexicon [29], or statistical approaches that assume the availability of a large dataset annotated with polarity or emotion labels [30].

Several supervised and unsupervised classifiers have been built to recognize emotional content in texts [31]. The SNoW architecture [32] is one of the most useful frameworks for text-based emotion detection. In the last decade, researchers have been focusing on sentiment extraction from texts of different genres, such as



Fig. 1. The hourglass of emotion.

product reviews [33], news [34], tweets [35], and essays [36], to name a few.

3.2. Audio visual emotion and sentiment analysis

In 1970, Ekman et al. [37] carried out extensive studies on facial expressions. Their research showed that universal facial expressions are able to provide sufficient clues to detect emotions. They used anger, sadness, surprise, fear, disgust, and joy as six basic emotion classes. Such basic affective categories are sufficient to describe most of the emotions expressed by facial expression. However, this list does not include the emotion expressed through facial expression by a person when he or she shows disrespect to someone; thus, a seventh basic emotion, contempt, was introduced by Matsumoto [38].

The Active Appearance Model [39,40] and Optical Flow-based techniques [41] are common approaches that use facial expression coding system (FACS) to understand facial expressions. Exploiting action units (AU) as features in well known classifier like k nearest neighbors, Bayesian networks, hidden Markov models (HMM), and artificial neural networks (ANN) [42] has helped many researchers to infer emotions from facial expression. The performance of several machine-learning algorithms for detecting emotions from facial expressions is presented in Table 1 (cited from Chen et al. [43]). All such systems, however, use different, manually-crafted corpora, which makes it impossible to perform a comparative evaluation of their performance.

To this end, recently Xu et al. [44] constructed a framework which takes color features from the superpixel of images and later a piece-wise linear transformation was used to learn the

Table 1

Performance of various learning algorithms for detecting emotions from facial images.

Method	Processing	Classification algorithm	Accuracy
Lanitis et al. [39]	Appearance Model	Distance-based	74%
Cohen et al. [45]	Appearance Model	Bayesian network	83%
Mase [46]	Optical flow	kNN	86%
Rosenblum et al. [47]	Optical flow	ANN	88%
Otsuka and Ohya [48]	2D FT of optical flow	HMM	93%
Yacoub and Davis [49]	Optical flow	Rule-based	95%
Essa and Pentland [50]	Optical flow	Distance-based	98%

emotional feature distribution. The framework is basically a novel feature learning framework from emotion labeled set of images.

Recent studies on speech-based emotion analysis [40,51–54] have focused on identifying several acoustic features such as fundamental frequency (pitch), intensity of utterance [43], bandwidth, and duration. The speaker-dependent approach gives much better results than the speaker-independent approach, as shown by the excellent results of Navas et al. [55], where about 98% accuracy was achieved by using the Gaussian mixture model (GMM) as a classifier, with prosodic, voice quality as well as Mel frequency cepstral coefficients (MFCC) employed as speech features.

When it comes to fusing audio-visual emotion recognition, two of the early works were done by De Silva et al. [56,57]. Both of these works showed that a bimodal system yielded a higher accuracy than any unimodal system. More recent research on audio-visual emotion recognition has been conducted at either feature level [58–60] or decision level [61–64]. Though there are plenty of research articles on audio-visual emotion recognition, only a few pieces of research works have been done on multimodal emotion or sentiment analysis using textual clues along with visual and audio modality. The works as described in [4,5,65] fused information from audio, visual and textual modalities to extract emotion and sentiment. Metallinou et al. [66,67] fused audio and textual modality for emotion recognition. Both these approaches relied on feature-level fusion. Wu et al. [68] fused audio and textual clues at decision level.

3.3. Multiple kernel learning

Several studies have reported that MKL outperforms the average kernel baselines. MKL is very similar to group LASSO which is a feature selection method where features are organized into groups. However, the choice of kernel coefficients can have a significant impact on the classification accuracy and efficiency of MKL [69].

For example, in Alzheimer's disease patients, different types of tests correspond to different modalities that can reveal varied aspects of the diagnosis. MR images may show only a slight hippocampal atrophy while the FDG-PET image may reveal increased hypometabolism suggestive of Alzheimer. In [70], MKL was used simultaneously for optimizing different modalities in Alzheimer's disease. However, in order to deal with co-morbidity with other diseases, they used the hinge loss function to penalize misclassified samples that did not scale well with the number of kernels. Adaptive MKL (AdaMKL) was proposed in [71] based on biconvex optimization and Gaussian kernels. Here, the objective function alternatively learns one component at a time while fixing the others, resulting in an increased computation time.

In [72], higher order kernels are used to enhance the learning of MKL. Here, block co-ordinate gradient optimization is used as it approximates the Hessian matrix of derivatives as a diagonal resulting in loss of information. MKL is also used in signal processing where grouping of features is useful to improve the interpretability of the learned parameters [73].

MKL was applied to a Polish opinion aggregator service that contained textual opinions of different products, but this study did not consider the hierarchical relation of different attributes of products [74]. Group-sensitive MKL for object recognition in images integrates a global kernel clustering method with MKL for sharing of group-sensitive information [75]. Hence, the two different kernels are used to group the training data and the kernels are aligned during optimization. They showed that their method outperformed baseline grouping strategies on the WikipediaMM dataset of real-world web images. The drawback of this method is that a looping strategy is used to relabel groups and may not reach the global optimum solution.

MKL was used to detect the presence of a large lump in images using a convolution kernel [76]. However, they only considered Gaussian features for the images. In [77], MKL was used to combine and re-weight multiple features by using structured latent variables during video event detection [77]. Here, two different types of kernels are used to group global features and segments in the test video that are similar to the training videos. While, the results on TRECVID dataset of video events outperformed baselines, the method requires tuning of parameters and assumes random initialization of latent variables.

Multimodal features were fused at different levels of fusion for the indexing of web data in [78]. The concept of kernel slack variables for each of the base kernels was used to classify YouTube videos in [79]. In order to select good features and discard bad features that may not be useful to the kernel, Liu et al. [80] used a beta prior distribution. Recently, MKL with Fourier transform on the Gaussian kernels has been applied to Alzheimer's Disease classification using both sMRI and fMRI images [81]. Researchers used L2 norm to enforce group sparsity constraints which were not robust on noisy datasets. Lastly, Online MKL shows good accuracy on object recognition tasks by extending online kernel learning to online MKL, however, the time complexity of the methods is dependent on the dataset [82].

4. Dataset used

In this section, we describe the datasets used in multimodal sentiment and emotion analysis experiments.

4.1. Multimodal sentiment analysis dataset

For our experiment, we use the dataset developed by Morency et al. [2]. They started collecting the videos from popular social media (e.g., YouTube) using several keywords to produce search results consisting of videos of either product reviews or recommendation. Some of these keywords are *my favorite products*, *non recommended perfumes*, *recommended movies* etc. A total of 80 videos were collected in this way. The dataset includes 15 male and 65 female speakers, with their age ranging approximately from 20–60 years.

The videos were converted to mp4 format with a standard size of 360 × 480. All videos were pre-processed to avoid the issues of introductory titles and multiple topics, and the length of the videos

Table 2
Utterances per emotion class.

Angry	Happy	Sad	Neutral	Total
1083	1630	1083	1683	5479

varied from 2 to 5 min. Many videos on YouTube contained an introductory sequence where a title was shown, sometimes accompanied with a visual animation. To address this issue, the first 30 seconds was removed from each video. Morency et al. [2] provided transcriptions with the videos. Each video was segmented into its utterances and each utterance was labeled by a sentiment, thanks to [2]. Because of the annotation scheme of the dataset, textual data was available for our experiment. On average each video has 6 utterances and each utterance is 5 seconds long. The dataset contains 498 utterances labeled either positive, negative or neutral. In our experiment we did not consider neutral labels, which led to the final dataset consisting of 448 utterances.

4.2. Multimodal emotion analysis dataset

The USC IEMOCAP database [6] was collected for the purposes of studying multimodal expressive dyadic interactions. This dataset contains 12 hours of video data split into 5 min of dyadic interaction between professional male and female actors. It was assumed that the interaction between the speakers are more affectively enriched than a speaker reading an emotional script. Each interaction session was split into spoken utterances. At least 3 annotators assigned the one emotion category, i.e., *happy*, *sad*, *neutral*, *angry*, *surprised*, *excited*, *frustration*, *disgust*, *fear* and *other* to each utterance. In this research work, we consider only the utterances with majority agreement (i.e., at least two out of three annotators labeled the same emotion) in the emotion classes of: Angry, Happy, Sad, and Neutral. Table 2 shows the per emotion class distribution.

5. Extracting features from visual data

Humans are known to express emotions through facial expression, to a great extent. As such, these expressions play a significant role in the identification of emotions in a multimodal stream. A facial expression analyzer automatically identifies emotional clues associated with facial expressions, and classifies these expressions to define sentiment categories and discriminate between them. We use positive and negative as sentiment classes in the classification problem. In the annotations provided with the YouTube dataset, each video was segmented into utterances and each of the utterances has the length of a few seconds. Every utterance was annotated as either 1, 0 and -1 , denoting positive, neutral and negative sentiment. Using a matlab code, we converted all videos in the dataset to image frames, after which we extracted facial features from each image frame. To extract facial characteristic points (FCPs) from the images, we used the facial recognition library CLM-Z [7]. From each image we extracted 68 FCPs; see examples in Table 3. The FCPs were used to construct facial features, which were defined as distances between FCPs; see examples in Table 4.

GAVAM[83] was also used to extract facial expression features from the face. Table 5 shows the extracted features from facial images. In our experiment we used the features extracted by CLM-Z along with the features extracted using GAVAM.

If a segment of a video has n number of images, then we extracted features from each image and take mean and standard deviation of those feature values in order to compute the final facial expression feature vector for an utterance.

Table 3
Some relevant facial characteristic points (out of the 68 facial characteristic points detected by CLM-Z).

Features	Description
48	Left eye
41	Right eye
43	Left eye inner corner
46	Left eye outer corner
47	Left eye lower line
44	Left eye upper line
40	Right eye inner corner
37	Right eye outer corner
42	Right eye lower line
38	Right eye upper line
23	Left eyebrow inner corner
25	Left eyebrow middle
27	Left eyebrow outer corner
22	Right eyebrow inner corner
20	Right eyebrow middle
18	Right eyebrow outer corner
52	Mouth top
58	Mouth bottom
55	Mouth Left corner
49	Mouth Right Corner
14	Middle of the Left Mouth Side
4	Middle of the Right Mouth Side

Table 4
Some important facial features used for the experiment.

Features
Distance between right eye and left eye.
Distance between the inner and outer corner of the left eye.
Distance between the upper and lower line of the left eye.
Distance between the left iris corner and right iris corner of the left eye.
Distance between the inner and outer corner of the right eye.
Distance between the upper and lower line of the right eye.
Distance between the left eyebrow inner and outer corner.
Distance between the right eyebrow inner and outer corner.
Distance between top of the mouth and bottom of the mouth.
Distance between left and right mouth corner.
Distance between the middle point of left and right mouth side.
Distance between Lower nose point and upper mouth point.

Table 5
Features extracted using GAVAM from the facial features.

Features
The time of occurrence of the particular frame in milliseconds.
The displacement of the face w.r.t X-axis. It is measured by the displacement of the normal to the frontal view of the face in the X-direction.
The displacement of the face w.r.t Y-axis.
The displacement of the face w.r.t Z-axis.
The angular displacement of the face w.r.t X-axis. It is measured by the angular displacement of the normal to the frontal view of the face with the X-axis.
The angular displacement of the face w.r.t Y-axis.
The angular displacement of the face w.r.t Z-axis.

6. Extracting features from audio data

We automatically extracted audio features from each annotated segment of the videos. Audio features were also extracted in 30Hz frame-rate and we used a sliding window of 100 ms. To compute the features we used the open source software openSMILE [84]. Specifically, this toolkit automatically extracts pitch and voice intensity. Z-standardization was used to perform voice normalization. Basically, voice normalization was performed and voice intensity was thresholded to identify samples with and without voice. The features extracted by openSMILE consist of several low-level descriptors (LLD) and their statistical functionals. Some of the functionals are *amplitude mean*, *arithmetic mean*, *root quadratic mean*, *standard deviation*, *flatness*, *skewness*, *kurtosis*, *quartiles*,

inter-quartile ranges, linear regression slope etc. Taking into account all functionals of each LLD, we obtained 6373 features. Some of the useful key LLD extracted by openSMILE are described below.

- Mel frequency cepstral coefficients – MFCC were calculated based on short time Fourier transform (STFT). First, log-amplitude of the magnitude spectrum was taken, and the process was followed by grouping and smoothing the fast Fourier transform (FFT) bins according to the perceptually motivated Mel-frequency scaling.
- Spectral Centroid – Spectral Centroid is the center of gravity of the magnitude spectrum of the STFT. Here, $M_i[n]$ denotes the magnitude of the Fourier transform at frequency bin n and frame i . The centroid is used to measure the spectral shape. A higher value of the centroid indicates brighter textures with greater frequency. The spectral centroid is calculated as

$$C_i = \frac{\sum_{n=0}^n n M_i[n]}{\sum_{n=0}^n M_i[n]}$$

- Spectral Flux – Spectral Flux is defined as the squared difference between the normalized magnitudes of successive windows:

$$F_i = \sum_{n=1}^n (N_t[n] - N_{t-1}[n])^2$$

where $N_t[n]$ and $N_{t-1}[n]$ are the normalized magnitudes of the Fourier transform at the current frame t and the previous frame $t - 1$, respectively. The spectral flux represents the amount of local spectral change.

- Beat histogram – It is a histogram showing the relative strength of different rhythmic periodicities in a signal. It is calculated as the auto-correlation of the RMS.
- Beat sum – This feature is measured as the sum of all entries in the beat histogram. It is a very good measure of the importance of regular beats in a signal.
- Strongest beat – It is defined as the strongest beat in a signal, in beats per minute, and it is found by identifying the strongest bin in the beat histogram.
- Pause duration – Pause duration is the percentage of time the speaker is silent in the audio segment.
- Pitch – It is computed by the standard deviation of the pitch level for a spoken segment.
- Voice Quality – Harmonics to noise ratio in the audio signal.
- PLP – The Perceptual Linear Predictive Coefficients of the audio segment were calculated using the openSMILE toolkit.

7. Extracting features from textual data

For feature extraction from textual data, we used a CNN. The trained CNN features were then fed into a SVM for classification. So, in particular we used CNN as trainable feature extractor and SVM as a classifier.

The intuition for building this hybrid classifier SVM-CNN is to combine the merits of each classifier and form a hybrid classifier to enhance accuracy. Recent studies [85] also show the use of CNN for feature extraction. In theory, the training process of CNN is similar to MLP as CNN is an extension of traditional MLP. MLP network is trained using a back-propagation algorithm which uses Empirical Risk Minimization. It tries to minimize the errors in training data. Once it finds the hyperplane, regardless of global or local optimum, the training process is stopped. This means that it does not try to improve the separation of the instances from the hyperplane. Wherein, SVM tries to minimize the generalization error on unseen data based on Structural Risk Minimization algorithm using a fixed probability distribution on training data. It therefore aims to maximize the distance between training instances and hyperplane,

so the margin area between two separate training classes is maximized. This separating hyperplane is a global optimum solution. So, SVM is more generalized than MLP which enhances the classification accuracy.

On the other hand, CNN automatically extracts key features from the training data. It grasps contextual local features from a sentence and after several convolution operations it finally forms a global feature vector out of those local features. CNN does not need the hand-crafted features used in a traditional supervised classifier. The hand-crafted features are difficult to compute and a good guess for encoding the features is always necessary in order to get satisfactory result. CNN uses a hierarchy of local features which are important to learn context. The hand-crafted features often ignore such a hierarchy of local features. Features extracted by CNN can therefore be used instead of hand-crafted features, as they carry more useful information.

The hybrid classifier SVM-CNN therefore inherits the merits from each classifier and should produce a better result.

The idea behind convolution is to take the dot product of a vector of k weights w_k also known as kernel vector with each k -gram in the sentence $s(t)$ to obtain another sequence of features $c(t) = (c_1(t), c_2(t), \dots, c_L(t))$.

$$c_j = w_k^T \cdot x_{i:i+k-1} \quad (1)$$

We then apply a max pooling operation over the feature map and take the maximum value $\hat{c}(t) = \max\{c(t)\}$ as the feature corresponding to this particular kernel vector. Similarly, varying kernel vectors and window sizes are used to obtain multiple features [86].

For each word $x_i(t)$ in the vocabulary, an d dimensional vector representation is given in a look up table that is learned from the data [87]. The vector representation of a sentence is hence a concatenation of vectors for individual words. Similarly we can have look up tables for other features. One might want to provide features other than words if these features are suspected to be helpful. The convolution kernels are then applied to word vectors instead of individual words.

We use these features to train higher layers of the CNN, to represent bigger groups of words in sentences. We denote the feature learned at hidden neuron h in layer l as F_h^l . Multiple features may be learned in parallel in the same CNN layer. The features learned in each layer are used to train the next layer

$$F^l = \sum_{h=1}^{n_h} w_k^h * F^{l-1} \quad (2)$$

where $*$ indicates convolution and w_k is a weight kernel for hidden neuron h and n_h is the total number of hidden neurons. The CNN sentence model preserves the order of words by adopting convolution kernels of gradually increasing sizes that span an increasing number of words and ultimately the entire sentence. Each word in a sentence was represented using word embedding and part-of-speech of that word. The details are as follows –

- **Word Embeddings** – We employ the publicly available word2vec vectors that were trained on 100 billion words from Google News. The vectors have dimensionality 300 trained using the continuous bag-of-words architecture [87]. Words not present in the set of pre-trained words are initialized randomly.
- **Part of Speech** – The part of speech of each word was also appended to the word's vector representation. As there are a total of 6 part of speech, so the length of part of speech vector was 6.

So, in the end a word was represented by a 306 dimensional vector.

Each sentence was wrapped to a window of 50 words to reduce the number of parameters and hence over-fitting the model. The CNN we developed in our experiment had two convolution layers,

a kernel size of 3 and 50 feature maps was used in the first convolution layer and a kernel size 2 and 100 feature maps in the second. It should be noted that the output of each convolution hidden layer is computed using a non-linear function (in our case we use tanh). Each convolution layer was followed by a max-pool layer. The max-pool size of the first and second max-pool layer was 2. The penultimate max-pool layer is followed by a fully connected layer with softmax output. We used 500 neurons in the full connected layer. The output layer corresponded to two neurons for each class of sentiments.

We used the output of the fully connected layer (layer 6) of the network as our feature vector. This feature vector was used in the final fusion process. So, in the fusion the 500 dimensional textual vector was used.

7.1. Other sentence-level textual features

We have ultimately fed the features extracted by CNN to the SVM and MKL. Motivated by the state of the art [88], we have decided to use other sentence-level features with the CNN extracted features. Below, we explain these features –

- **Commonsense knowledge features** – Commonsense knowledge features consist of concepts are represented by means of AffectiveSpace [89]. In particular, concepts extracted from text through the semantic parser are encoded as 100-dimensional real-valued vectors and then aggregated into a single vector representing the sentence by coordinate-wise summation:

$$x_i = \sum_{j=1}^N x_{ij},$$

where x_i is the i -th coordinate of the sentence's feature vector, $i = 1, \dots, 100$; x_{ij} is the i th coordinate of its j th concept's vector, and N is the number of concepts in the sentence (extracted by means of our concept parser [90]).

- **Sentic feature** – The polarity scores of each concept extracted from the sentence were obtained from SenticNet and summed up to produce a single scalar feature.
- **Part-of-speech feature** – This feature is defined by the number of adjectives, adverbs and nouns in the sentence, which give three distinct features.
- **Modification feature** – This is a single binary feature. For each sentence, we obtained its dependency tree from the dependency parser. This tree was analyzed to determine whether there is any word modified by a noun, adjective, or adverb. The modification feature is set to 1 in case of any modification relation in the sentence; 0 otherwise.
- **Negation feature** – Similarly, the negation feature is a single binary feature determined by the presence of any negation in the sentence. It is important because the negation can invert the polarity of the sentence.

8. Experimental results

For the experiment, we removed all neutral classes resulting in the final dataset of 448 utterances. Of these, 247 were negative and 201 were positive. In this section, we describe the experimental results of the unimodal and multimodal frameworks. For each experiment, we carried out 10-fold cross validation.

8.1. Extracting sentiment from visual modality

To extract sentiment from only visual modality we used SVM classifier with a polykernel. Features were extracted using the method explained in Section 5. Table 6 shows the results for each class –{positive and negative}.

Table 6

Confusion matrix for the visual modality (SVM classifier).

Actual classification	Predicted classification			
	Negative	Positive	Precision	Recall
Negative	197	50	76.40%	79.80%
Positive	61	140	73.70%	69.70%

Table 7

Confusion matrix for the audio modality (SVM classifier).

Actual classification	Predicted classification			
	Negative	Positive	Precision	Recall
Negative	208	38	73.20%	84.60%
Positive	76	125	76.70%	62.20%

Table 8

Confusion matrix for the textual modality (CNN classifier).

Actual classification	Predicted classification			
	Negative	Positive	Precision	Recall
Negative	210	36	78.65%	85.36%
Positive	57	143	79.88%	71.50%

Clearly, the recall is lower for positive samples. This means many negative instances were labelled as positive. Below, we show some features which took major role to confuse the classifier.

- The large change in distance of FCPs on eyelid from lower eyebrow.
- Small change between the two corners of the mouth (F_{49} and F_{55} as shown in Fig. 2).

We compared the performance of SVM with other classifiers like Multilayer Perceptron (MLP) and Extreme Learning Machine (ELM) [91]. SVM was found to produce best performance results. On visual modality, the best state-of-the-art result on this dataset was obtained by Rosas et al. [5] where they got 67.31% accuracy. In terms of accuracy our method has outperformed their result by achieving 75.22% accuracy.

8.2. Extracting sentiment from audio modality

For each utterance, we extracted the features as stated in Section 6 and formed a feature vector which was then fed to SVM. Table 7 shows that for the positive class, the classifier obtained relatively lower recall than for the visual modality obtained. Rosas et al. [5] obtained 64.85% accuracy on audio modality. Conversely, a 74.49% accuracy was obtained using the proposed method, outperforming the accuracy of the state-of-the-art-model [5]. For 1 utterance in the dataset, there is no audio data. This resulted in 447 utterances in the final dataset for this experiment.

8.3. Extracting sentiment from textual modality

As we described in Section 7, deep Convolutional Network (CNN) was used to extract features from textual modality and a SVM classifier was then employed on those features to identify sentiment. We call this hybrid classifier CNN-SVM. Comparing the performance of CNN-SVM with other supervised classifiers, we found it to offer the best classification results (Table 8). In this experiment, our method also outperformed the state-of-the-art accuracy achieved by Rosas et al. [5]. For 2 utterances, no text data was available in the dataset. So, the final dataset for this experiment consists of 446 utterances out of which 246 are negative and 200 are positive.




Fig. 2. A sample of facial characteristic points extracted by CLM-Z.

Table 9

Confusion matrix for the Audio-Visual Modality (MKL Classifier).

Actual classification	Predicted classification			
	Negative	Positive	Precision	Recall
Negative	214	32	82.90%	87.00%
Positive	44	156	83.00%	78.00%

The results shown in Table 8 were obtained when the utterances in the dataset were translated from Spanish to English. Without this translation process we obtained a much lower accuracy of 68.56%. Another experimental study showed that while using CNN-SVM produced a 79.14% accuracy, an accuracy of only 75.50% was achieved using CNN.

8.4. Feature-level fusion of audio, visual and textual modalities

After extracting features from all modalities, we merged them to form a long feature vector. That feature vector was then fed to MKL for the classification task. We tested several polynomial kernels of different degree and RBF kernels having different gamma values as base kernels in MKL. We compared the performance of SPG-GMKL (Spectral Projected Gradient-Generalized Multiple Kernel Learning)[92] and Simple-MKL in the classification task and found that SPG-GMKL outperformed Simple-MKL with a 1.3% relative error reduction rate. Based on the cross validation performance, the best set of kernels and their corresponding parameters were chosen. Finally, we chose a configuration with 8 kernels: 5 RBF with gamma from 0.01 to 0.05 and 3 polynomial with powers 2, 3, 4.

Table 9 shows the results of the audio-visual feature-level fusion. Clearly, the performance in terms of both precision and recall increased when these two modalities are fused.

Among the unimodal classifiers, textual modality was found to provide the most accurate classification result. We observed the same fact when textual features were fused with audio and visual

Table 10

Confusion matrix for the Audio-Textual Modality (SPG-GMKL Classifier).

Actual classification	Predicted classification			
	Negative	Positive	Precision	Recall
Negative	217	29	83.46%	88.21%
Positive	43	157	84.40%	78.50%

Table 11

Confusion matrix for the Visual-Textual Modality (MKL Classifier).

Actual classification	Predicted classification			
	Negative	Positive	Precision	Recall
Negative	221	25	84.03%	89.83%
Positive	42	158	86.33%	79.00%

Table 12

Confusion matrix for the Audio-Visual-Textual Modality (SPG-GMKL Classifier).

Actual classification	Predicted classification			
	Negative	Positive	Precision	Recall
Negative	227	19	86.64%	92.27%
Positive	35	165	89.67%	82.50%

modalities. Both the audio-textual (Table 10) and visual-textual (Table 11) framework outperformed the audio-visual framework. According to the experimental results, visual-textual modality performed best.

Table 13 shows the results when all three modalities were fused producing a 87.89% accuracy. Clearly, this accuracy is higher than the best state-of-the-art framework, which obtained a 74.09% accuracy. The fundamental reason for our method outperforming the state-of-the-art method is the extraction of salient features from each modality before fusing those features using MKL.

Table 13

Results and Comparison of Unimodal experiment and Multimodal Feature-Level Fusion (Accuracy).

	Perez-Rosa et al. [5]	Our Method
Audio Modality	64.85%	74.49%
Visual Modality	67.31%	75.22%
Textual Modality	70.94%	79.14%
Visual and text-based features	72.39%	84.97%
Visual and audio-based features	68.86%	82.95%
Audio and text-based features	72.88%	83.85%
Fusing all three modalities	74.09%	87.89%

Table 14

Results and Comparison of Unimodal experiment and Multimodal Feature-Level Fusion (Accuracy): Feature Selection was carried out.

	Perez-Rosa et al. [5]	Our Method
Audio Modality	64.85%	74.22%
Visual Modality	67.31%	76.38%
Textual Modality	70.94%	79.77%
Visual and text-based features	72.39%	85.46%
Visual and audio-based features	68.86%	83.69%
Audio and text-based features	72.88%	84.12%
Fusing all three modalities	74.09%	88.60%

8.5. Feature selection

In order to see whether a reduced optimal feature subset can produce a better result than using all features, we conducted a cyclic Correlation-based Feature Subset Selection (CFS) using the training set of each fold. The main idea of CFS is that useful feature subsets should contain features that are highly correlated with the target class while being uncorrelated with each other. However, superior results were obtained when we used all features. This signifies that some relevant features were excluded by CFS. We then employed Principal Component Analysis (PCA) for feature selection to rank all features according to their importance in classification. To measure whether top K features selected PCA can produce better accuracy, we fed the top K features to the classifier. However, even worse accuracy was obtained than when using CFS based feature selection. When we took the combination of top K features from that ranking and CFS-based selected features and employed the classifier on them, we observed the best accuracy. To set the value of K, an exhaustive search was made and finally we found that K=300 gave the best result. This evaluation was carried out for each experiment stated in Sections 8.1, 8.2, and 8.4.

For our audio, visual and textual fusion experiment using CFS and PCA, a total 437 features were selected out of which 305 features were textual, 74 were visual and 58 were from audio modality. This proves the fact that textual features were the most important for trimodal sentiment analysis thanks to CNN feature extractor. Table 14 shows the comparative evaluation using feature selection method.

8.6. Feature-level fusion for multimodal emotion recognition

Besides doing the experiment on multimodal sentiment analysis dataset, we also carried out an extensive experiment on multimodal emotion analysis dataset as described in Section 4.2. We followed the same method as applied for the sentiment analysis dataset. However, instead of taking it as a binary classification task, we considered it as a 4-way classification. This dataset already provides the facial points detected by the markers and we only used those facial points in our study. CLM-Z was not able to detect faces in most of the facial images as the images in this dataset are small and of low resolution. Using a similar feature selection algorithm as described in Section 8.5, a total of 693 features were selected,

of which 85 features were textual, 239 were audio and 369 were from visual modality.

In Table 15 we see that both precision and recall of the Happy class is higher. However, Angry and Sad classes are very tough to distinguish from the textual clues. One of the possible reasons is both of these classes are negative emotions and many words are commonly used to express both of the emotions. On the other hand, the classifier was confused and often classified Neutral with Happy and Anger. Interestingly, it classifies Sad and Neutral classes well.

In the case of Audio modality (Table 16) we observe better accuracy than textual modality for Sad and Neutral classes. However, for Happy and Angry, the performance decreased. The confusion matrix shows the classifier performed poorly when distinguishing Angry from Happy. Clearly, audio features are unable to effectively classify these based on extracted features. However, the classifier performs very well to discriminate between the classes of Sad and Anger. Overall identification accuracy of the Neutral emotion has also increased. But Happy and Neutral emotions are still very hard to classify effectively by Audio classifier alone.

Visual modality produced the best accuracy (Table 17) when compared to other two modalities. The similar trend has been observed as textual modality. Angry and Sad faces are hard to classify using visual clues. However, Angry and Happy, Happy and Sad faces can be effectively classified. Neutral classes were also separated accurately in respect to other classes.

When we fuse the modalities using the feature-level fusion strategy (Table 18) as stated in Section 8.4, as expected higher accuracy was obtained than with unimodal classifiers. Although the identification accuracy has been improved for every emotion, the confusion between a Sad and Angry face is still higher. Neutral and Sad emotions are also more difficult to classify.

The comparison with the state-of-the-art model in terms of weighted accuracy shows that the proposed method performs significantly better. Comparing the weighted accuracy (WA) with the state of the art, the proposed method obtained 3.75% higher accuracy. However, for Anger emotion class, an approximately 3% lower accuracy was achieved.

8.7. Decision-level fusion

In this section, we describe different frameworks that we developed for the decision-level fusion. Clearly, the motivation for developing these frameworks is to perform the fusion process in less time. The fusion frameworks were developed according to the architecture as shown in Fig. 3. Each of the experiments stated below were processed through the feature selection algorithm stated in Section 8.5.

Each block M_i denotes a modality. As the architecture shows, modality M_1 and M_2 are fused using feature-level fusion and then at last stage are fused with another modality M_3 using decision-level fusion. For feature-level fusion of M_1 and M_2 , we used SPG-GMKL. The decision-level algorithm is described below -

In decision-level fusion, we obtained the feature vectors from the above-mentioned methods but used separate classifier for each modality instead of concatenating feature vectors as in feature-level fusion. The output of each classifier was treated as a classification score. In particular, from each classifier we obtained a probability score for each sentiment class. In our case, as there are two sentiment classes, we obtained 2 probability scores from each modality. Let, q_1^{12} and q_2^{12} are the class probabilities resulted from the feature-level fusion of M_1 and M_2 . On the other hand let, q_1^3 and q_2^3 are the class probabilities of modality M_3 . We then form a feature vector by concatenating these class probabilities.

We also used sentic patterns [94] to obtain the sentiment label for each text. If the result by sentic patterns for a sentence

Table 15

Confusion matrix for the Textual Modality (SVM Classifier, Feature selection carried out).

Actual classification	Predicted classification				Precision	Recall
	Angry	Happy	Sad	Neutral		
Angry	650	82	165	186	55.13%	60.01%
Happy	193	957	149	331	68.40%	58.71%
Sad	139	87	619	238	55.51%	57.15%
Neutral	197	273	182	1031	57.72%	61.25%

Table 16

Confusion matrix for the Audio Modality (SVM Classifier, Feature selection carried out).

Actual classification	Predicted classification				Precision	Recall
	Angry	Happy	Sad	Neutral		
Angry	648	137	89	209	61.53%	59.83%
Happy	159	926	123	422	65.21%	56.81%
Sad	84	152	658	189	63.08%	60.75%
Neutral	162	205	173	1143	58.22%	67.91%

Table 17

Confusion matrix for the Visual Modality (SVM Classifier, Feature selection carried out).

Actual classification	Predicted classification				Precision	Recall
	Angry	Happy	Sad	Neutral		
Angry	710	83	116	174	66.17%	65.55%
Happy	102	1034	148	346	72.76%	63.43%
Sad	123	83	726	151	63.18%	67.03%
Neutral	138	221	159	1165	63.45%	69.22%

Table 18

Confusion matrix for the Audio-Visual-Textual Modality (SPG-GMKL Classifier, Feature selection carried out).

Actual classification	Predicted classification				Precision	Recall
	Angry	Happy	Sad	Neutral		
Angry	821	79	93	90	69.16%	75.80%
Happy	119	1217	92	202	80.11%	74.67%
Sad	93	82	782	126	67.24%	72.20%
Neutral	154	141	196	1192	73.99%	70.82%

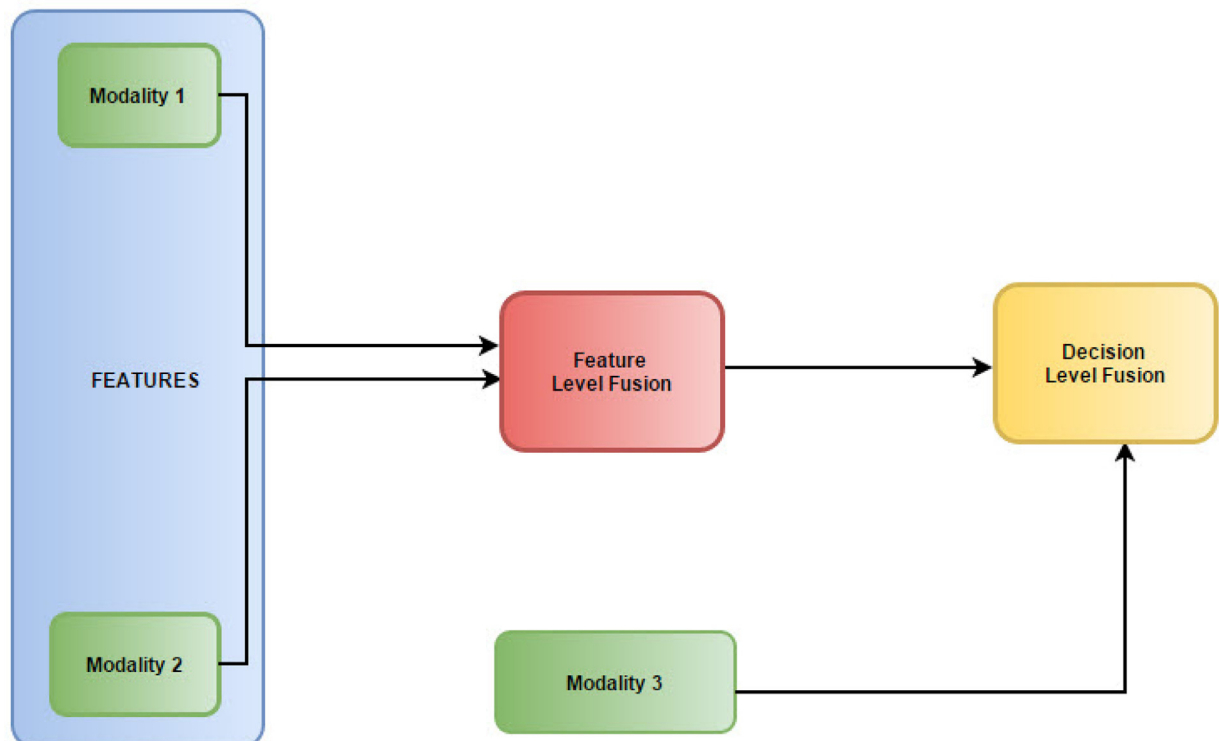
**Fig. 3.** Decision-level fusion framework.

Table 19

Comparison with the state of the art [93] on IEMOCAP dataset.

	Rozic et al. [93]	Proposed Method
Anger	78.10%	75.80%
Happy	69.20%	74.67%
Sad	67.10%	72.20%
Neutral	63.00%	70.82%
WA	69.50%	73.25%

Table 20

Decision-level fusion accuracy.

M_1	M_2	M_3	Sentic Patterns	Accuracy
Visual	Audio	Textual	No	73.31%
Visual	Audio	Textual	Yes	78.30%
Visual	Textual	Audio	No	76.62%
Audio	Textual	Visual	No	72.50%

Table 21

Decision-Level Fusion Accuracy for Multimodal Emotion Analysis.

M_1	M_2	M_3	Weighted Accuracy
Visual	Audio	Textual	64.20%
Visual	Textual	Audio	62.75%
Audio	Textual	Visual	61.22%

is "positive" then we included 1 in the feature vector, otherwise 0 was included in the feature vector. So, the final feature vector looks like this - $[q_1^{12}, q_2^{12}, q_3^3, q_2^3, \text{sentic}]$ where $\text{sentic} = 1$ if the output of sentic patterns is positive otherwise we set $\text{sentic} = 0$. We then employed SVM on this feature vector in order to obtain the final polarity label.

The best accuracy was obtained when we early fused visual and audio modalities. However, when we fuse all the modalities without carrying out the early fusion, the obtained accuracy was lower. Table 20 shows the decision-level accuracy in detail.

8.7.1. Decision-level fusion for multimodal emotion detection

Like decision-level fusion for multimodal sentiment analysis, similar method was applied for multimodal emotion analysis as well (Table 21).

Similarly as we saw in the sentiment analysis experiment, the configuration yielding best accuracy was obtained using M_1 , M_2 and M_3 as Visual, Audio and Textual, respectively.

Table 19 shows the detail result of decision-level fusion experiment on IEMOCAP dataset. It should be noted that sentic patterns cannot be used in this experiment as it is specific to sentiment analysis.

9. Speeding up the computational time: The role of ELM

9.1. Extreme learning machine

The ELM approach [95] was introduced to overcome some issues in back-propagation network [96] training, specifically; potentially slow convergence rates, the critical tuning of optimization parameters, and the presence of local minima that call for multi-start and re-training strategies. The ELM learning problem settings require a training set, X , of N labeled pairs, where (\mathbf{x}_i, y_i) , where $\mathbf{x}_i \in \mathcal{R}^m$ is the i th input vector and $y_i \in \mathcal{R}$ is the associate expected 'target' value; using a scalar output implies that the network has one output unit, without loss of generality.

The input layer has m neurons and connects to the 'hidden' layer (having N_h neurons) through a set of weights $\{\mathbf{w}_j \in \mathcal{R}^m; j = 1, \dots, N_h\}$. The j th hidden neuron embeds a bias term, \hat{b}_j , and a nonlinear 'activation' function, $\varphi(\cdot)$; thus the neuron's response

to an input stimulus, \mathbf{x} , is:

$$a_j(\mathbf{x}) = \varphi(\mathbf{w}_j \cdot \mathbf{x} + \hat{b}_j) \quad (3)$$

Note that (3) can be further generalized as a wider class of functions [97] but for the subsequent analysis this aspect is not relevant. A vector of weighted links, $\mathbf{w}_j \in \mathcal{R}^{N_h}$, connects hidden neurons to the output neuron without any bias [98]. The overall output function, $f(\mathbf{x})$, of the network is:

$$f(\mathbf{x}) = \sum_{j=1}^{N_h} \mathbf{w}_j a_j(\mathbf{x}) \quad (4)$$

It is convenient to define an 'activation matrix', \mathbf{H} , such that the entry $\{h_{ij} \in \mathbf{H}; i = 1, \dots, N; j = 1, \dots, N_h\}$ is the activation value of the j th hidden neuron for the i th input pattern. The \mathbf{H} matrix is:

$$\mathbf{H} \equiv \begin{bmatrix} \varphi(\mathbf{w}_1 \cdot \mathbf{x}_1 + \hat{b}_1) & \cdots & \varphi(\mathbf{w}_{N_h} \cdot \mathbf{x}_1 + \hat{b}_{N_h}) \\ \vdots & \ddots & \vdots \\ \varphi(\mathbf{w}_1 \cdot \mathbf{x}_N + \hat{b}_1) & \cdots & \varphi(\mathbf{w}_{N_h} \cdot \mathbf{x}_N + \hat{b}_{N_h}) \end{bmatrix} \quad (5)$$

In the ELM model, the quantities $\{\mathbf{w}_j, \hat{b}_j\}$ in (3) are set randomly and are not subject to any adjustment, and the quantities $\{\mathbf{w}_j, \hat{b}\}$ in (4) are the only degrees of freedom. The training problem reduces to the minimization of the convex cost:

$$\min_{\{\mathbf{w}, \hat{b}\}} \|\mathbf{H}\mathbf{w} - \mathbf{y}\|^2 \quad (6)$$

A matrix pseudo-inversion yields the unique L_2 solution, as proven in [95]:

$$\mathbf{w} = \mathbf{H}^+ \mathbf{y} \quad (7)$$

The simple and efficient procedure to train an ELM therefore involves the following steps:

1. Randomly set the input weights \mathbf{w}_i and bias \hat{b}_i for each hidden neuron;
2. Compute the activation matrix, \mathbf{H} , as per (5);
3. Compute the output weights by solving a pseudo-inverse problem as per (7).

Despite the apparent simplicity of the ELM approach, the crucial result is that even random weights in the hidden layer endow a network with a notable representation ability [95]. Moreover, the theory derived in [99] proves that regularization strategies can further improve its generalization performance. As a result, the cost function (6) is augmented by an L_2 regularization factor as follows:

$$\min_{\mathbf{w}} \{\|\mathbf{H}\mathbf{w} - \mathbf{y}\|^2 + \lambda \|\mathbf{w}\|^2\} \quad (8)$$

9.2. Experiment and comparison with SVM

The experimental results in Table 22 shows ELM and SVM offering equivalent performance in terms of accuracy. While for multimodal sentiment analysis SVM outperformed ELM with a sharp 1.23% accuracy margin, on the emotion analysis dataset their performance difference is not significant. On the IEMOCAP dataset, ELM showed better accuracy for text based emotion detection. Importantly, for the purposes of feature-level fusion, we used a multiple kernel variant of the ELM algorithm namely multiple-kernel extreme learning machine (MK-ELM) [100]. As with SPG-GMKL for feature-level fusion (Section 8.4), the same set of kernels was used for MK-ELM.

However, ELM edges SVM out by a big margin when it comes to computational time, i.e., training time of feature-level fusion (see Table 23).

SPG-GMKL outperformed SVM for the feature-level fusion task by 2.7%.

Table 22

Accuracy Comparison between SVM and ELM (A=Audio, V=Video, T=Textual, UWA=Un-weighted Average)

Dataset	A		V		T		A+V+T (UWA)	
	SVM	ELM	SVM	ELM	SVM	ELM	SPG-GMKL	MK-ELM
YouTube	74.22%	73.81%	76.38%	76.24%	79.77%	78.36%	88.60%	87.33%
IEMOCAP	61.32%	60.85%	66.30%	64.74%	59.28%	59.87%	73.37	72.68%

Table 23

Computational Time comparison between SVM and ELM.

	YouTube Dataset	IEMOCAP dataset
SPG-GMKL	1926 seconds	4389 seconds
MK-ELM	584 seconds	2791 seconds

10. Conclusion

In this work, a novel multimodal affective data analysis framework is proposed. It includes the extraction of salient features, development of unimodal classifiers, building feature- and decision-level fusion frameworks. The deep CNN-SVM-based textual sentiment analysis component is found to be the key element for outperforming the state-of-the-art model's accuracy. MKL has played a significant role in the fusion experiment. The novel decision-level fusion architecture is also an important contribution of this paper. In the case of the decision-level fusion experiment, the coupling of sentic patterns to determine the weight of textual modality has enriched the performance of the multimodal sentiment analysis framework considerably.

Interestingly, a lower accuracy was obtained for the emotion recognition task, which may indicate that extracting emotions from video may be more difficult than inferring polarity. While text is the most important factor for determining polarity, the visual modality shows the best performance for emotion analysis. The most interesting part of this paper is that a common multimodal affect data analysis framework is well capable of extracting emotion and sentiment from different datasets.

Future work will focus on extracting more relevant features via visual modality. Specifically, deep 3D CNNs will be employed for automatic feature extraction from videos. A feature selection method will be used to select only the best features in order to ensure both scalability and stability of the framework. Consequently, we will strive to improve the decision-level fusion process using a cognitive inspired fusion engine. In order to realize our ambitious goal of developing a novel real-time system for multimodal sentiment analysis, the time complexities of the methods need to be consistently reduced. Hence, another aspect of our future work will be to effectively analyze and appropriately address the system's time complexity requirements in order to create a better, more time efficient and reliable multimodal sentiment analysis engine.

References

- [1] E. Cambria, Affective computing and sentiment analysis, *IEEE Intel. Syst.* 31 (2) (2016) 102–107.
- [2] L.-P. Morency, R. Mihalcea, P. Doshi, Towards multimodal sentiment analysis: Harvesting opinions from the web, in: *Proceedings of the 13th International Conference on Multimodal Interfaces*, ACM, 2011, pp. 169–176.
- [3] E. Cambria, N. Howard, J. Hsu, A. Hussain, Sentic blending: Scalable multimodal fusion for continuous interpretation of semantics and sentics, in: *Proceedings of the IEEE SSCI*, Singapore, 2013, pp. 108–117.
- [4] M. Wollmer, F. Weninger, T. Knaup, B. Schuller, C. Sun, K. Sagae, L.-P. Morency, Youtube movie reviews: Sentiment analysis in an audio-visual context, *IEEE Intell. Syst.* 28 (3) (2013) 46–53.
- [5] V. Rosas, R. Mihalcea, L.-P. Morency, Multimodal sentiment analysis of spanish online videos, *IEEE Intel. Syst.* 28 (3) (2013) 0038–45.
- [6] C. Busso, M. Bulut, C.-C. Lee, A. Kazemzadeh, E. Mower, S. Kim, J.N. Chang, S. Lee, S.S. Narayanan, Iemocap: Interactive emotional dyadic motion capture database, *Lang. Resources Eval.* 42 (4) (2008) 335–359.
- [7] T. Baltrusaitis, P. Robinson, L. Morency, 3d constrained local model for rigid and non-rigid facial tracking, in: *Proceedings of the 2012 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2012, pp. 2610–2617.
- [8] H. Qi, X. Wang, S.S. Iyengar, K. Chakrabarty, Multisensor data fusion in distributed sensor networks using mobile agents, in: *Proceedings of 5th International Conference on Information Fusion*, 2001, pp. 11–16.
- [9] S. Poria, I. Chaturvedi, E. Cambria, A. Hussain, Convolutional MKL based multimodal emotion recognition and sentiment analysis, in: *Proceedings of the ICDM, Barcelona*, 2016, pp. 439–448.
- [10] S. Poria, E. Cambria, R. Bajpai, A. Hussain, A review of affective computing: From unimodal analysis to multimodal fusion, *Inf. Fusion* 37 (2017) 98–125.
- [11] E. Cambria, H. Wang, B. White, Guest editorial: Big social data analysis, *Knowl.-Based Syst.* 69 (2014) 1–2.
- [12] B. Pang, L. Lee, S. Vaithyanathan, Thumbs up?: Sentiment classification using machine learning techniques, in: *Proceedings of the EMNLP, ACL*, 2002, pp. 79–86.
- [13] R. Socher, A. Perelygin, J.Y. Wu, J. Chuang, C.D. Manning, A.Y. Ng, C. Potts, Recursive deep models for semantic compositionality over a sentiment treebank, in: *Proceedings of EMNLP*, 1631, 2013, pp. 1642–1654.
- [14] H. Yu, V. Hatzivassiloglou, Towards answering opinion questions: Separating facts from opinions and identifying the polarity of opinion sentences, in: *Proceedings of the EMNLP, ACL*, 2003, pp. 129–136.
- [15] P. Melville, W. Gryc, R.D. Lawrence, Sentiment analysis of blogs by combining lexical knowledge with text classification, in: *Proceedings of the ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, ACM, 2009, pp. 1275–1284.
- [16] A. Zadeh, R. Zellers, E. Pincus, L.-P. Morency, Multimodal sentiment intensity analysis in videos: Facial gestures and verbal messages, *IEEE Intel. Syst.* 31 (6) (2016) 82–88.
- [17] X. Hu, J. Tang, H. Gao, H. Liu, Unsupervised sentiment analysis with emotional signals, in: *Proceedings of the WWW*, 2013, pp. 607–618.
- [18] A. Gangemi, V. Presutti, D. Reforgiato Recupero, Frame-based detection of opinion holders and topics: A model and a tool, *IEEE Comput. Intel. Mag.* 9 (1) (2014) 20–30, doi:10.1109/MCI.2013.2291688.
- [19] E. Cambria, A. Hussain, *Sentic Computing: A Common-Sense-Based Framework for Concept-Level Sentiment Analysis*, Springer, Cham, Switzerland, 2015.
- [20] E. Cambria, S. Poria, R. Bajpai, B. Schuller, SenticNet 4: A semantic resource for sentiment analysis based on conceptual primitives, in: *Proceedings of the COLING*, 2016, pp. 2666–2677.
- [21] G. Qiu, B. Liu, J. Bu, C. Chen, Expanding domain sentiment lexicon through double propagation, in: *Proceedings of the IJCAI*, vol.9, 2009, pp. 1199–1204.
- [22] H. Kanayama, T. Nasukawa, Fully automatic lexicon expansion for domain-oriented sentiment analysis, in: *Proceedings of the EMNLP, ACL*, 2006, pp. 355–363.
- [23] J. Blitzer, M. Dredze, F. Pereira, et al., Biographies, bollywood, boom-boxes and blenders: Domain adaptation for sentiment classification, in: *Proceedings of the ACL* 2007, vol.7, 2007, pp. 440–447.
- [24] S.J. Pan, X. Ni, J.-T. Sun, Q. Yang, Z. Chen, Cross-domain sentiment classification via spectral feature alignment, in: *Proceedings of the WWW, ACM*, 2010, pp. 751–760.
- [25] D. Bollegala, D. Weir, J. Carroll, Cross-domain sentiment classification using a sentiment sensitive thesaurus, *IEEE Trans. Knowl. Data Eng.* 25 (8) (2013) 1719–1731.
- [26] C. Strapparava, A. Valitutti, Wordnet affect: an affective extension of wordnet, in: *Proceedings of the LREC*, 4, 2004, pp. 1083–1086.
- [27] C.O. Alm, D. Roth, R. Sproat, Emotions from text: machine learning for text-based emotion prediction, in: *Proceedings of the EMNLP, Association for Computational Linguistics*, 2005, pp. 579–586.
- [28] E. Cambria, A. Livingstone, A. Hussain, The hourglass of emotions, in: *Cognitive Behavioural Systems*, Springer, 2012, pp. 144–157.
- [29] G. Mishne, Experiments with mood classification in blog posts, in: *Proceedings of ACM SIGIR 2005 Workshop on Stylistic Analysis of Text for Information Access*, 19, 2005.
- [30] L. Oneto, F. Bisio, E. Cambria, D. Anguita, Statistical learning theory and ELM for big social data analysis, *IEEE Comput. Intel. Mag.* 11 (3) (2016) 45–55.
- [31] C. Yang, K.H.-Y. Lin, H.-H. Chen, Building emotion lexicon from weblog corpora, in: *Proceedings of the 45th Annual Meeting of the ACL on Interactive Poster and Demonstration Sessions*, Association for Computational Linguistics, 2007, pp. 133–136.

- [32] F.-R. Chaumartin, Upar7: A knowledge-based system for headline sentiment tagging, in: Proceedings of the 4th International Workshop on Semantic Evaluations, Association for Computational Linguistics, 2007, pp. 422–425.
- [33] S. Poria, E. Cambria, A. Gelbukh, Aspect extraction for opinion mining with a deep convolutional neural network, *Knowl.-Based Syst.* 108 (2016) 42–49.
- [34] X. Li, H. Xie, L. Chen, J. Wang, X. Deng, News impact on stock price return via sentiment analysis, *Knowl.-Based Syst.* 69 (2014) 14–23.
- [35] P. Chikersal, S. Poria, E. Cambria, A. Gelbukh, C.E. Siong, Modelling public sentiment in twitter: using linguistic patterns to enhance supervised learning, in: Computational Linguistics and Intelligent Text Processing, Springer, 2015, pp. 49–65.
- [36] S. Poria, A. Gelbukh, B. Agarwal, E. Cambria, N. Howard, Common sense knowledge based personality recognition from text, in: *Advances in Soft Computing and Its Applications*, Springer, pp. 484–496.
- [37] P. Ekman, Universal facial expressions of emotion, *Culture and Personality: Contemporary Readings/Chicago*, 1974.
- [38] D. Matsumoto, More evidence for the universality of a contempt expression, *Motiv. Emotion* 16 (4) (1992) 363–368.
- [39] A. Lanitis, C.J. Taylor, T.F. Cootes, A unified approach to coding and interpreting face images, in: Proceedings of the Fifth International Conference on Computer Vision, IEEE, 1995, pp. 368–373.
- [40] D. Datzu, L. Rothkrantz, Semantic audio-visual data fusion for automatic emotion recognition, in: Proceedings of the Euromedia, 2008.
- [41] M. Kenji, Recognition of facial expression from optical flow, *IEICE Trans. Inform. Syst.* 74 (10) (1991) 3474–3483.
- [42] N. Ueki, S. Morishima, H. Yamada, H. Harashima, Expression analysis/synthesis system based on emotion space constructed by multilayered neural network, *Syst. Comput. Japan* 25 (13) (1994) 95–107.
- [43] L.S.-H. Chen, Joint processing of audio-visual information for the recognition of emotional expressions in human-computer interaction, 2000 Ph.D. thesis.
- [44] M. Xu, B. Ni, J. Tang, S. Yan, Image re-emotionalizing, in: *The Era of Interactive Media*, Springer, 2013, pp. 3–14.
- [45] I. Cohen, N. Sebe, A. Garg, L.S. Chen, T.S. Huang, Facial expression recognition from video sequences: temporal and static modeling, *Comput. Vis. Image Underst.* 91 (1) (2003) 160–187.
- [46] M. Mansoorizadeh, N.M. Charkari, Multimodal information fusion application to human emotion recognition from face and speech, *Multim. Tools Appl.* 49 (2) (2010) 277–297.
- [47] M. Rosenblum, Y. Yacoob, L.S. Davis, Human expression recognition from motion using a radial basis function network architecture, *IEEE Trans. Neural Netw.* 7 (5) (1996) 1121–1138.
- [48] T. Otsuka, J. Ohya, A study of transformation of facial expressions based on expression recognition from temporal image sequences, Technical Report, Institute of Electronic, Information, and Communications Engineers (IEICE), 1997.
- [49] Y. Yacoob, L.S. Davis, Recognizing human facial expressions from long image sequences using optical flow, *IEEE Trans. Pattern Anal. Mach. Intell.* 18 (6) (1996) 636–642.
- [50] I.A. Essa, A.P. Pentland, Coding, analysis, interpretation, and recognition of facial expressions, *IEEE Trans. Pattern Anal. Mach. Intell.* 19 (7) (1997) 757–763.
- [51] I.R. Murray, J.L. Arnott, Toward the simulation of emotion in synthetic speech: A review of the literature on human vocal emotion, *J. Acoust. Soc. Am.* 93 (2) (1993) 1097–1108.
- [52] R. Cowie, E. Douglas-Cowie, Automatic statistical analysis of the signal and prosodic signs of emotion in speech, in: Proceedings of the Fourth International Conference on Spoken Language, 3, IEEE, 1996, pp. 1989–1992.
- [53] F. Dellaert, T. Polzin, A. Waibel, Recognizing emotion in speech, in: Proceedings of the Fourth International Conference on Spoken Language, 1996. ICSLP 96, 3, IEEE, 1996, pp. 1970–1973.
- [54] T. Johnstone, Emotional speech elicited using computer games, in: Proceedings of the Fourth International Conference on Spoken Language, 1996. ICSLP 96, 3, IEEE, 1996, pp. 1985–1988.
- [55] E. Navas, I. Hernaez, I. Luengo, An objective and subjective study of the role of semantics and prosodic features in building corpora for emotional tts, *IEEE Trans. Audio, Speech, Lang. Process.* 14 (4) (2006) 1117–1127.
- [56] L.C. De Silva, T. Miyasato, R. Nakatsu, Facial emotion recognition using multi-modal information, in: Proceedings of 1997 International Conference on Information, Communications and Signal Processing, 1, IEEE, 1997, pp. 397–401.
- [57] L.S. Chen, T.S. Huang, T. Miyasato, R. Nakatsu, Multimodal human emotion/expression recognition, in: Proceedings of the Third IEEE International Conference on Automatic Face and Gesture Recognition, 1998, IEEE, 1998, pp. 366–371.
- [58] Y. Wang, L. Guan, Recognizing human emotional state from audiovisual signals, *IEEE Trans. Multim.* 10 (5) (2008) 936–946.
- [59] D. Datzu, L.J. Rothkrantz, Emotion recognition using bimodal data fusion, in: Proceedings of the 12th International Conference on Computer Systems and Technologies, ACM, 2011, pp. 122–128.
- [60] L. Kessous, G. Castellano, G. Caridakis, Multimodal emotion recognition in speech-based interaction using facial expression, body gesture and acoustic analysis, *J. Multimodal User Interf.* 3 (1–2) (2010) 33–48.
- [61] B. Schuller, Recognizing affect from linguistic information in 3d continuous space, *IEEE Trans. Affective Comput.* 2 (4) (2011) 192–205.
- [62] M. Rashid, S. Abu-Bakar, M. Mokji, Human emotion recognition from videos using spatio-temporal and audio features, *Vis. Comput.* 29 (12) (2013) 1269–1275.
- [63] M. Glodek, S. Reuter, M. Schels, K. Dietmayer, F. Schwenker, Kalman filter based classifier fusion for affective state recognition, in: *Multiple Classifier Systems*, Springer, 2013, pp. 85–94.
- [64] S. Hommel, A. Rabie, U. Handmann, Attention and emotion based adaption of dialog systems, in: *Intelligent Systems: Models and Applications*, Springer, 2013, pp. 215–235.
- [65] V. Rozgic, S. Ananthakrishnan, S. Saleem, R. Kumar, R. Prasad, Speech language & multimedia technol., raytheon bbn technol., cambridge, ma, usa, in: Proceedings of the 2012 Asia-Pacific Signal & Information Processing Association Annual Summit and Conference (APSIPA ASC), IEEE, 2012, pp. 1–4.
- [66] A. Metallinou, S. Lee, S. Narayanan, Audio-visual emotion recognition using gaussian mixture models for face and voice, in: Proceedings of the Tenth IEEE International Symposium on Multimedia, IEEE, 2008, pp. 250–257.
- [67] F. Eyben, M. Wöllmer, A. Graves, B. Schuller, E. Douglas-Cowie, R. Cowie, On-line emotion recognition in a 3-d activation-valence-time continuum using acoustic and linguistic cues, *J. Multimodal User Interf.* 3 (1–2) (2010) 7–19.
- [68] C.-H. Wu, W.-B. Liang, Emotion recognition of affective speech based on multiple classifiers using acoustic-prosodic information and semantic labels, *IEEE Trans. Affect. Comput.* 2 (1) (2011) 10–21.
- [69] S. Bucak, R. Jin, A. Jain, Multiple kernel learning for visual object recognition: a review, *IEEE Trans. Pattern Anal. Mach. Intell.* 36 (7) (2014) 1354–1369.
- [70] C. Hinrichs, V. Singh, G. Xu, S. Johnson, Mkl for robust multi-modality ad classification, *Med. Image Comput. Comput. Assist. Interv.* 5762 (2009) 786–794.
- [71] Z. Zhang, Z.-N. Li, M. Drew, Adamkl: A novel biconvex multiple kernel learning approach, in: Proceedings of the 2010 20th International Conference on Pattern Recognition (ICPR), 2010, pp. 2126–2129.
- [72] S. Wang, S. Jiang, Q. Huang, Q. Tian, Multiple kernel learning with high order kernels, in: Proceedings of the International Conference on Pattern Recognition, 2010, pp. 2138–2141.
- [73] N. Subrahmanya, Y. Shin, Sparse multiple kernel learning for signal processing applications, *IEEE Trans. Pattern Anal. Mach. Intell.* 32 (5) (2010) 788–798.
- [74] A. Wawer, Mining opinion attributes from texts using multiple kernel learning, in: Proceedings of 2011 IEEE 11th International Conference on Data Mining Workshops (ICDMW), 2011, pp. 123–128.
- [75] J. Yang, Y. Tian, L.-Y. Duan, T. Huang, W. Gao, Group-sensitive multiple kernel learning for object recognition, *IEEE Trans. Image Process.* 21 (5) (2012) 2838–2852.
- [76] S. Nilufar, N. Ray, H. Zhang, Object detection with dog scale-space: A multiple kernel learning approach, *IEEE Trans. Image Process.* 21 (8) (2012) 3744–3756.
- [77] A. Vahdat, K. Cannons, G. Mori, S. Oh, I. Kim, Compositional models for video event detection: a multiple kernel learning latent variable approach, in: Proceedings of 2013 IEEE International Conference on Computer Vision (ICCV), 2013, pp. 1185–1192.
- [78] U. Niaz, B. Merialdo, Fusion methods for multi-modal indexing of web data, in: Proceedings of the 2013 14th International Workshop on Image Analysis for Multimedia Interactive Services (WIAMIS), 2013, pp. 1–4.
- [79] X. Xu, I. Tsang, D. Xu, Soft margin multiple kernel learning, *IEEE Trans. Neural Netw. Learn. Syst.* 24 (5) (2013) 749–761.
- [80] B. Ni, T. Li, P. Moulin, Beta process multiple kernel learning, in: Proceedings of the 2014 IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2014, pp. 963–970.
- [81] F. Liu, L. Zhou, C. Shen, J. Yin, Multiple kernel learning in the primal for multimodal alzheimers disease classification, *IEEE J. Biomedical Health Inform.* 18 (3) (2014) 984–990.
- [82] H. Xia, S. Hoi, R. Jin, P. Zhao, Online multiple kernel similarity learning for visual search, *IEEE Trans. Pattern Anal. Mach. Intell.* 36 (3) (2014) 536–549.
- [83] J.M. Saragih, S. Lucey, J.F. Cohn, Face alignment through subspace constrained mean-shifts, in: Proceedings of the 2009 IEEE 12th International Conference on Computer Vision, IEEE, 2009, pp. 1034–1041.
- [84] F. Eyben, M. Wöllmer, B. Schuller, Opensmile: the munich versatile and fast open-source audio feature extractor, in: Proceedings of the International Conference on Multimedia, ACM, 2010, pp. 1459–1462.
- [85] A.S. Razavian, H. Azizpour, J. Sullivan, S. Carlsson, Cnn features off-the-shelf: an astounding baseline for recognition, in: Proceedings of the 2014 IEEE Conference on Computer Vision and Pattern Recognition Workshops (CVPRW), IEEE, 2014, pp. 512–519.
- [86] N. Kalchbrenner, E. Grefenstette, P. Blunsom, A convolutional neural network for modelling sentences, *CoRR abs/1404.2188* (2014).
- [87] T. Mikolov, K. Chen, G. Corrado, J. Dean, Efficient estimation of word representations in vector space, (2013) arXiv preprint arXiv: 1301.3781.
- [88] S. Poria, E. Cambria, A. Gelbukh, Deep convolutional neural network textual features and multiple kernel learning for utterance-level multimodal sentiment analysis, in: Proceedings of EMNLP, 2015, pp. 2539–2544.
- [89] E. Cambria, J. Fu, F. Bisio, S. Poria, AffectiveSpace 2: Enabling affective intuition for concept-level sentiment analysis, in: Proceedings of the AAAI, Austin, 2015, pp. 508–514.
- [90] D. Rajagopal, E. Cambria, D. Olsher, K. Kwok, A graph-based approach to commonsense concept extraction and semantic similarity detection, in: Proceedings of the WWW, Rio De Janeiro, 2013, pp. 565–570.
- [91] G.-B. Huang, E. Cambria, K.-A. Toh, B. Widrow, Z. Xu, New trends of learning in computational intelligence, *IEEE Comput. Intel. Mag.* 10 (2) (2015) 16–17.

- [92] A. Jain, S. Vishwanathan, M. Varma, Spf-gmkl: generalized multiple kernel learning with a million kernels, in: Proceedings of the 18th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, ACM, 2012, pp. 750–758.
- [93] V. Rozgic, S. Ananthakrishnan, S. Saleem, R. Kumar, R. Prasad, Ensemble of svm trees for multimodal emotion recognition, in: Proceedings of the 2012 Asia-Pacific Signal & Information Processing Association Annual Summit and Conference (APSIPA ASC), IEEE, 2012, pp. 1–4.
- [94] S. Poria, E. Cambria, A. Gelbukh, F. Bisio, A. Hussain, Sentiment data flow analysis by means of dynamic linguistic patterns, *IEEE Comput. Intel. Mag.* 10 (4) (2015) 26–36.
- [95] G.-B. Huang, D.H. Wang, Y. Lan, Extreme learning machines: a survey, *Int. J. Mach. Learn. Cybern.* 2 (2) (2011) 107–122.
- [96] S. Ridella, S. Rovetta, R. Zunino, Circular backpropagation networks for classification, *IEEE Trans. Neural Netw.* 8 (1) (1997) 84–97.
- [97] G.-B. Huang, L. Chen, C.-K. Siew, Universal approximation using incremental constructive feedforward networks with random hidden nodes, *IEEE Trans. Neural Netw.* 17 (4) (2006) 879–892.
- [98] G.-B. Huang, An insight into extreme learning machines: Random neurons, random features and kernels, *Cognitive Comput.* (2014), doi:10.1007/s12559-014-9255-2.
- [99] G.-B. Huang, H. Zhou, X. Ding, R. Zhang, Extreme learning machine for regression and multiclass classification, *IEEE Trans. Syst. Man, Cybern. Part B: Cybern.* 42 (2) (2012) 513–529.
- [100] X. Liu, L. Wang, G.-B. Huang, J. Zhang, J. Yin, Multiple kernel extreme learning machine, *Neurocomputing* 149, Part A (2015) 253–264.



Soujanya Poria received his BEng in Computer Science from Jadavpur University, India in 2013. He then joined Nanyang Technological University as a research engineer in the School of Electrical and Electronics Engineering and, later in 2015, he joined NTU Temasek Labs, where he is conducting research on sentiment analysis in multiple domains and different modalities. Since February 2014, Soujanya has also started his PhD studies at the University of Stirling (Computing Science and Mathematics). His research areas include natural language processing, opinion mining, cognitive science and multimodal sentiment analysis. In 2013, Soujanya received the best undergraduate thesis and researcher award from Jadavpur University.

He was awarded Gold Plated Silver medal from the University and Tata Consultancy Service for his final year project during his undergraduate course. He is also a fellow of the Brain Sciences Foundation and a program committee member of SENTIRE, the IEEE ICDM workshop series on sentiment analysis.



Haiyun Peng received his Bachelor of Engineering in automation from Wuhan University in 2013. After that, he obtained his Master of Science in Signal processing from Nanyang Technological University. He is currently a PhD student under supervision of Erik Cambria in the School of Computer Engineering in Nanyang Technological University. His main research interests are concept-level natural language processing and multi-modal sentiment analysis, both in English and Chinese language.



Amir Hussain obtained his BEng (with the highest 1st Class Honors) and PhD (in novel neural network architectures and algorithms) from the University of Strathclyde in Glasgow, Scotland, UK, in 1992 and 1997 respectively. He is currently a Professor of Computing Science, and founding Director of the Cognitive Signal Image and Control Processing Research (COSIPRA) Laboratory at the University of Stirling in Scotland, UK. His research interests are inter-disciplinary and industry focussed, and include multi-modal cognitive and sentic computing techniques and applications. He has published over 270 papers, including over a dozen books and 80 journal papers. He is the founding Editor-in-Chief of the journals: *Cognitive Computation* (Springer Neuroscience, USA), and *Big Data Analytics* (BioMed Central), and Chief-Editor of the Springer Book Series on Socio-Affective Computing, and Springer Briefs on Cognitive Computation. He is an Associate Editor of the IEEE Transactions on Neural Networks and Learning Systems, a member of several Technical Committees of the IEEE Computational Intelligence Society (CIS), founding publications co-Chair of the IINNS Big Data Section and its annual INNS Conference on Big Data, and Chapter Chair of the IEEE UK and RI Industry Applications Society.



Newton Howard's passion for science and technology began during his childhood. He pursued his interests in his studies and in 2000 while a graduate member of the Department of Mathematical Sciences at the University of Oxford, he proposed the Theory of Intention Awareness (IA). In 2002, he received a second doctoral degree in cognitive informatics and mathematics from the prestigious La Sorbonne in France. In 2007 he was awarded the habilitation a diriger des recherches (HDR) for his leading work on the Physics of Cognition (PoC) and its applications to complex medical, economical, and security equilibriums. Recently in 2014 he received his doctorate of philosophy from the University of Oxford specifically focusing on "The Brain Code" for work in neurodegenerative diseases. His work has made a significant impact on the design of command and control systems as well as information exchange systems used at tactical, operational and strategic levels. As the creator of IA, Dr. Howard was able to develop operational systems for military and law enforcement projects. These utilize an intent-centric approach to inform decision-making and ensure secure information sharing. His work has brought him into various academic and government projects of significant magnitude, which focus on science and the technological transfer to industry. While Dr. Howard's career formed in military scientific research, in 2002 he founded the Center for Advanced Defense Studies (CADS) a leading Washington, D.C, national security group. Currently, Dr. Howard serves as the Director of the Board. He also is a national security advisor to several U.S. Government organizations.



Erik Cambria received his BEng and MEng with honors in Electronic Engineering from the University of Genoa in 2005 and 2008, respectively. In 2012, he was awarded his PhD in Computing Science and Mathematics following the completion of an EPSRC project in collaboration with MIT Media Lab, which was selected as impact case study by the University of Stirling for the UK Research Excellence Framework (REF2014). After two long-term research visits at HP Labs India and Microsoft Research Asia, he worked as Lead Investigator in NUS Cognitive Science Programme till 2014. Today, Dr Cambria is an Assistant Professor at NTU School of Computer Science and Engineering, a Research Fellow at NTU Temasek Labs, and an Adjunct Scientist at A*STAR IHPC. His current affiliations also include Rolls-Royce@NTU, MIT Synthetic Intelligence Lab, and the Brain Sciences Foundation. He is Associate Editor of Elsevier KBS and IPM, IEEE CIM, Springer AIRE, Cognitive Computation, and Editor of the IEEE IS Department on Affective Computing and Sentiment Analysis.